A robot that acts to know

Day 1 — The challenge, probability, and free energy

Course built on the doctoral thesis of Noé Zapata

Monday 30 November

Montevideo, 30 November – 4 December 2026

The challenge

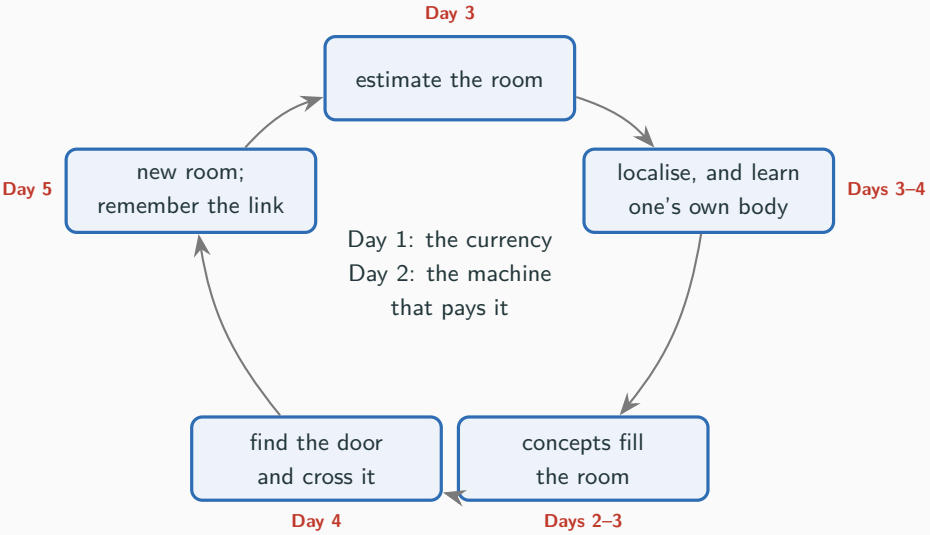
A robot is put down in an apartment it has never seen

[video: the full cycle, Webots / real apartment]

room estimated → localised → objects appear → door crossed → second room

Watch for one thing: nothing in this video was programmed as a sequence of steps.

The cycle we will build this week

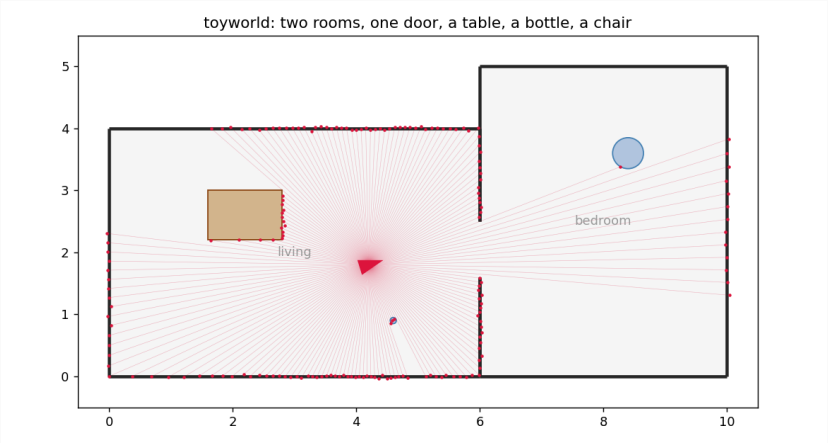


Every stage is a belief

stage	what is unknown (latent)	what supplies the evidence	day
estimate the room	walls, corners, the polygon	LiDAR returns	3
localise	the robot's pose in that room	LiDAR against the believed room	3
learn the body	wheel scale, gyro bias, sensor mounts	the localiser's corrections	4
concepts	which objects exist, where, what shape	camera masks, LiDAR	2–3
door	where it is, whether it is open	the room's boundary, free space	4
new room + memory	a second polygon; how rooms connect	all of the above, again	5

One generative model over all of it, and one quantity that every stage minimises: perception lowers it by changing beliefs, action by changing the data.

The toy version you will build, one stage a day



Two rooms, one door, three objects, a LiDAR, and a robot whose wheels lie a little. Pure `numpy`; it runs on your laptop.

How this course is built

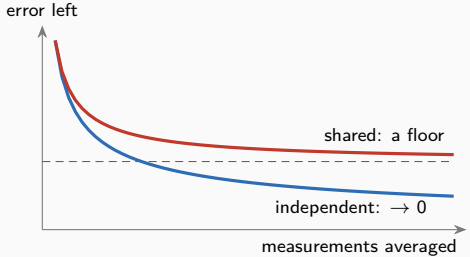
Each session runs through three steps:

1. the **theory**, derived rather than asserted;
2. the **code**, one stage of the cycle, in the lab;
3. the **measurements** that showed when the equations were wrong.

The third step is not optional. Most of what the thesis learned, it learned there.

The thread through all five days

Averaging removes independent error.
It does not remove shared error.

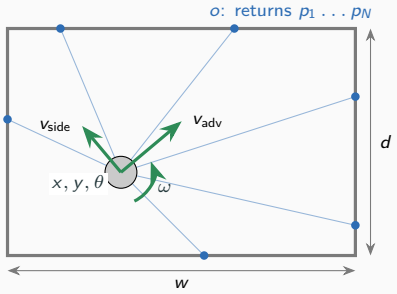


The grammar of uncertainty

What the robot is uncertain about

deterministic	$y = f(x)$: one outcome
random	$y = f(x) + \varepsilon$: many outcomes
discrete	a door is open *or* closed
continuous	a wall is 4.013 . . . m away

Everything the robot does not know is a **random variable**, and what it believes about it is a **probability distribution**: a table for discrete variables, a density for continuous ones.



s: pose and room o: scan a: command

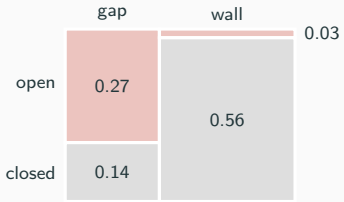
$$p(s, o, a)$$

The joint distribution: everything there is to know

A two-variable world. D: is the door open? O: does the LiDAR see a gap where the door is?

	$O =$ gap	$O =$ wall
$D =$ open	0.27	0.03
$D =$ closed	0.14	0.56

the four numbers add up to 1



The same **joint** $p(D, O)$, drawn so that *area is probability*. Everything this week is three operations on it: summing, slicing, and multiplying back.

Marginalisation: sum out what you do not care about

	gap	wall	$p(D)$
open	0.27	0.03	0.30
closed	0.14	0.56	0.70
$p(O)$	0.41	0.59	1

$$p(o) = \sum_d p(d, o) \qquad p(o) = \int p(s, o)\, \mathrm{d}s$$

- Sum each row or column and write the result *in the margin*: that is where the name comes from.
- A marginal answers a question about one variable while staying honest about every value the other one might take.
- The integral over s is the expensive one. Remember it.

Conditioning: keep the slice you observed, renormalise

	gap	wall
open	0.27	0.03
closed	0.14	0.56
$p(O)$	0.41	0.59

$$p(\text{open} \mid \text{gap}) = \frac{0.27}{0.41} = 0.66$$

$$p(s \mid o) = \frac{p(s, o)}{p(o)}$$

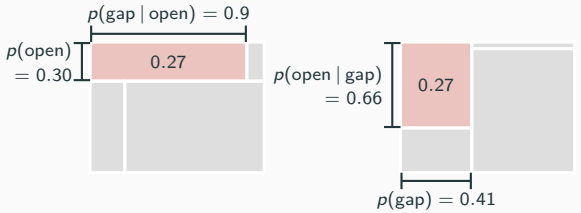


Before looking, open with probability 0.30; after a gap, 0.66; after a wall, $0.03/0.59 = 0.05$.
Conditioning is how evidence enters a belief.

Factorisation: the product rule

Turn the conditioning formula around:

$$p(s, o) = \underbrace{p(o \mid s)\, p(s)}_{\text{first } s\text{, then } o \text{ given } s} = \underbrace{p(s \mid o)\, p(o)}_{\text{first } o\text{, then } s \text{ given } o}$$



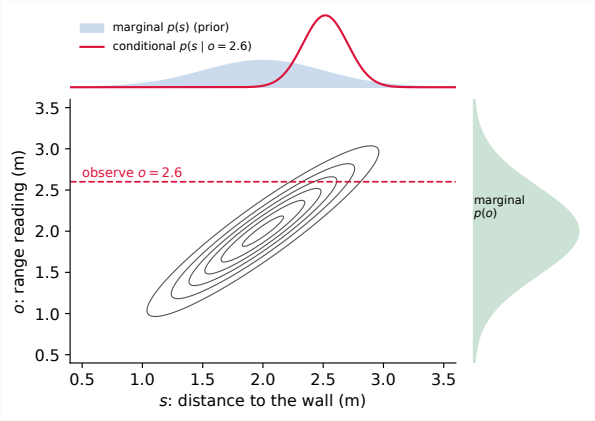
Two different rectangles, the same area: the same cell of the joint.

Set the two factorisations equal and divide by $p(o)$:

$$p(s \mid o) = \frac{p(o \mid s)\, p(s)}{p(o)}$$

Bayes' rule is not an extra axiom. It is the product rule written twice.

The same operations on continuous variables



s: distance to a wall, prior $\mathcal{N}(2, 0.5^2)$.
o: its range reading, $o \mid s \sim \mathcal{N}(s, 0.2^2)$.

- The **joint** is the tilted ellipse: readings track distances.
- A **marginal** is the shadow the joint casts on one axis.
- A **conditional** is one horizontal cut through it, rescaled to area 1.

Reading $o = 2.6$ moves the belief from 2.00 ± 0.50 to 2.52 ± 0.19 m.

source: figures/make_day1.py, block 6

For Gaussians, both operations are linear algebra

Split a Gaussian over (s, o) into blocks, written two ways:

$$\text{covariance } \Sigma = \begin{bmatrix} \Sigma_{ss} & \Sigma_{so} \\ \Sigma_{os} & \Sigma_{oo} \end{bmatrix} \qquad \text{precision } \Lambda = \Sigma^{-1} = \begin{bmatrix} \Lambda_{ss} & \Lambda_{so} \\ \Lambda_{os} & \Lambda_{oo} \end{bmatrix}$$

	in covariance form	in precision form
marginalise out o	keep the block: Σ_{ss}	Schur complement: $\Lambda_{ss} - \Lambda_{so}\Lambda_{oo}^{-1}\Lambda_{os}$
condition on o	Schur complement: $\Sigma_{ss} - \Sigma_{so}\Sigma_{oo}^{-1}\Sigma_{os}$	keep the block: Λ_{ss}

- Each form makes one operation free and the other costly. The slide before: $0.25 - 0.25^2/0.29 = 0.035 = 0.19^2$.
- The robot keeps its beliefs in precision form. When it forgets an old pose (Days 2–5), it is marginalising, and it pays a Schur complement.

The chain rule: N variables, N factors

Apply the product rule again and again, peeling off one variable each time:

$$
\begin{aligned}
p(A, B) &= p(A \mid B)\, p(B) \\
p(A, B, C) &= p(A \mid B, C)\, p(B \mid C)\, p(C) \\
p(A, B, C, D) &= p(A \mid B, C, D)\, p(B \mid C, D)\, p(C \mid D)\, p(D)
\end{aligned}
$$

In general, for any ordering π of the variables:

$$
p(X_1, \ldots, X_N) = \underbrace{p(X_{\pi(1)})\, p(X_{\pi(2)} \mid X_{\pi(1)}) \cdots p(X_{\pi(N)} \mid X_{\pi(1)}, \ldots, X_{\pi(N-1)})}_{N \text{ factors}}
$$

- No assumption was made. This is exact for every distribution.
- But the last factor depends on everything: nothing has been gained yet.

One joint, $N!$ factorisations

From now on the names we use all week: S the **state** (is the door open?), O the **observation** (does the LiDAR see a gap?), A the **action** (does the robot turn its LiDAR towards the door, or away?).

$$
\begin{aligned}
p(S, O, A) &= p(S)\, p(O \mid S)\, p(A \mid S, O) \\
&= p(S)\, p(A \mid S)\, p(O \mid S, A) \\
&= p(O)\, p(S \mid O)\, p(A \mid S, O) \\
&= p(O)\, p(A \mid O)\, p(S \mid O, A) \\
&= p(A)\, p(S \mid A)\, p(O \mid S, A) \\
&= p(A)\, p(O \mid A)\, p(S \mid O, A)
\end{aligned}
$$

- 3 variables give $3! = 6$ orderings, each with 3 factors.
- N variables give $N!$. A robot tracked over 100 steps has hundreds of variables.
- All of them are exactly equal. No amount of data can tell them apart.

So which one do we write down?

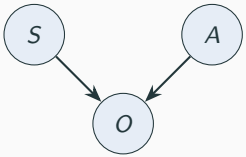
Choose the order in which the world produces things

Causal order: causes before effects.

$$p(S)\,p(A \mid S)\,p(O \mid S, A)$$

The robot cannot see whether the door is open, so where it points its LiDAR does not depend on it; what it observes depends on both:

$$p(A \mid S) = p(A)$$



Anti-causal order: effects first.

$$p(O)\,p(A \mid O)\,p(S \mid O, A)$$

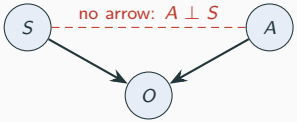
Once O is known, S and A are linked through it (a gap means little if the LiDAR looked away), so nothing simplifies. Every factor keeps all its arguments.

Numbers needed to specify it (binary variables):

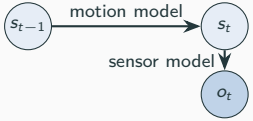
any order, nothing dropped	$2^3 - 1 = 7$
causal, using $A \perp S$	$1 + 1 + 4 = 6$
N variables, any order	$2^N - 1$
N variables, $\leq k$ causes each	$\leq N\,2^k$

Why the causal order wins

1. Sparse. The arrows you leave out are the independences you know.



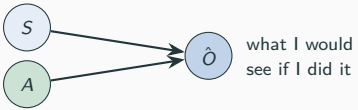
2. Each factor is physics. Nobody can write down $p(O)$, every possible scan.



3. Modular. Swap the LiDAR for a camera and see what must be re-learnt:

causal	$p(S)$	$p(A)$	$p(O \mid S, A)$
anti-causal	$p(O)$	$p(A \mid O)$	$p(S \mid O, A)$

4. It runs forwards. Sampling in order predicts what acting would show:



Three things you can build from one joint

Perception only

$p(s \mid o)$



A classifier, a detector.

Control only

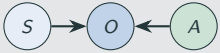
$p(a \mid o)$



PID, a hand-written controller, a learnt policy: the factor $p(A \mid O)$ of the anti-causal order, on its own.

The generative model

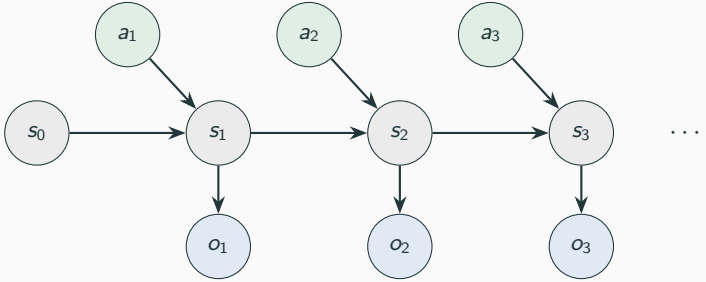
$p(s)\, p(a)\, p(o \mid s, a)$



Explains, predicts, and chooses by inference. This is what we build.

	acts, in the world it was tuned in	says what it expects to see	notices it is surprised	knows what it does not know	survives a sensor swap
controller, policy $p(a \mid o)$	✓	×	×	×	×
the joint	✓	✓	✓	✓	✓

The robot's joint, in causal order

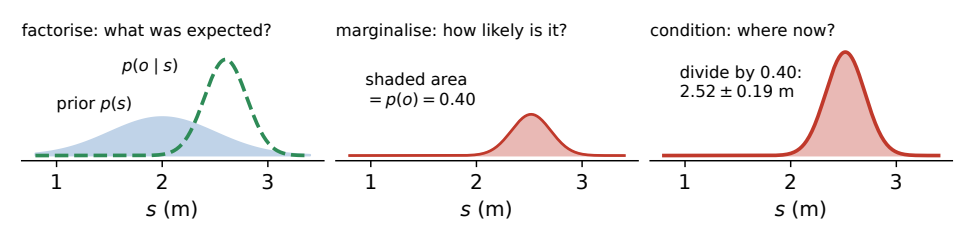


$$p(s_{0:T}, o_{1:T}, a_{1:T}) = \underbrace{p(s_0)}_{\text{prior}} \prod_{t=1}^{T} \underbrace{p(a_t)}_{\text{policy}} \underbrace{p(s_t \mid s_{t-1}, a_t)}_{\text{motion model}} \underbrace{p(o_t \mid s_t)}_{\text{sensor model}}$$

$(3T+1)!$ orderings are available, and we write one: each factor touches at most three variables and is a piece of physics. This is the **generative model**; tomorrow we divide it among the processes that actually compute it.

Bayes' rule is the three operations in a row

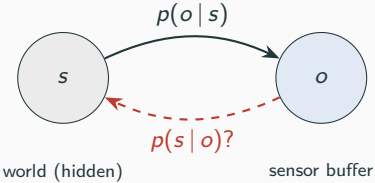
$$\underbrace{p(s \mid o)}_{\text{posterior}} = \frac{\overbrace{p(o \mid s)}^{\text{likelihood}}\ \overbrace{p(s)}^{\text{prior}}}{\underbrace{p(o)}_{\text{evidence}}}$$



The shaded area is where exact inference dies; the rest of today is how not to compute it.

Inference behind a sensory boundary

The agent never sees the state



The **generative model**:

$$p(s, o) = p(o \mid s)\, p(s)$$

- $p(o \mid s)$ is the agent's theory of its own sensor: which LiDAR returns a robot at this pose, in this room, would produce.
- $p(s)$ is what it brought with it: plausible room shapes, typical table sizes.

Inference runs the arrow backwards:

$$p(s \mid o) = \frac{p(o \mid s)\, p(s)}{p(o)}$$

source: active_inference.tex l.181–212

Where exact inference dies

$$p(o) = \int p(o \mid s)\, p(s)\, \mathrm{d}s$$

- The probability of *this scan* under *every* room and furniture configuration the robot considers possible, each weighted by how likely it is.
- No closed form and no feasible quadrature. This is not a quirk of one sensor: it is the generic situation.
- Every practical method from here on is a way of not computing it.

A word about "state"

s is a latent variable of the *agent's model*, not the physical truth, which is never accessible. That distinction is what gives the question "is this belief wrong?" its meaning.

source: active_inference.tex l.200–220

Variational free energy

Do not compute the posterior; approximate it

Choose a tractable family, and a member $q(s)$ of it (the **belief**). Move it towards the true posterior, measuring closeness by

$$D_{\mathrm{KL}}[q(s) \,\|\, p(s \mid o)] = \int q(s) \ln \frac{q(s)}{p(s \mid o)} \, \mathrm{d}s \;\geq 0.$$

Do not compute the posterior; approximate it

Choose a tractable family, and a member $q(s)$ of it (the **belief**). Move it towards the true posterior, measuring closeness by

$$D_{\mathrm{KL}}[q(s) \,\|\, p(s \mid o)] = \int q(s) \ln \frac{q(s)}{p(s \mid o)} \,\mathrm{d}s \;\geq 0.$$

Problem: it contains $p(s \mid o)$, the very thing we cannot have.

Do not compute the posterior; approximate it

Choose a tractable family, and a member $q(s)$ of it (the **belief**). Move it towards the true posterior, measuring closeness by

$$D_{\mathrm{KL}}[q(s) \,\|\, p(s \mid o)] = \int q(s) \ln \frac{q(s)}{p(s \mid o)} \,\mathrm{d}s \;\geq 0.$$

Problem: it contains $p(s \mid o)$, the very thing we cannot have.

Substitute Bayes' rule and separate out the part that does not depend on q:

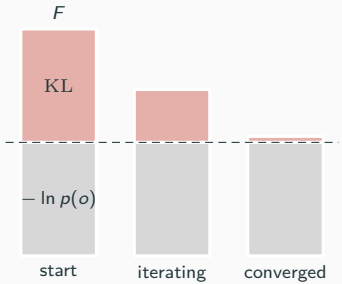
$$D_{\mathrm{KL}}[q \,\|\, p(s \mid o)] = \underbrace{\int q(s) \ln \frac{q(s)}{p(o, s)} \,\mathrm{d}s}_{F[q,o]} \;+\; \ln p(o)$$

F needs only the belief and the joint $p(o \mid s)p(s)$ the agent already wrote down.

source: active_inference.tex l.222–250

The identity the framework rests on

$$F[q, o] = D_{\mathrm{KL}}[q(s) \,\|\, p(s \mid o)] - \ln p(o)$$



dashed: no belief can move it

Same minimiser. Only the KL slab can shrink, so lowering the computable F lowers the uncomputable divergence.

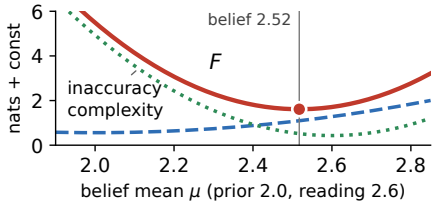
A bound on surprise. $\mathrm{KL} \geq 0$, so the bar never falls below the dashed line: $F \geq -\ln p(o)$.

If you know machine learning: $-F$ is the **ELBO**.

source: active_inference.tex l.252–268

Complexity minus accuracy

$$F[q, o] = \underbrace{D_{\mathrm{KL}}[q(s) \,\|\, p(s)]}_{\text{complexity}} - \underbrace{\mathbb{E}_q[\ln p(o \mid s)]}_{\text{accuracy}}$$



Accuracy alone would adopt any explanation that fits the current scan.

Complexity charges for the distance travelled from the prior: Occam's razor, written as a divergence.

The belief settles where the two slopes cancel.

One trade, stated once, in nats, for every quantity the agent believes.

source: active_inference.tex l.288–308

Worked example: the Kalman update is free-energy minimisation

Take the prior $s \sim \mathcal{N}(m, p)$, an observation $y \sim \mathcal{N}(s, r)$, and the belief $q = \mathcal{N}(\mu, \Sigma)$. Every term is a Gaussian expectation:

$$F(\mu, \Sigma) = \frac{(\mu - m)^2 + \Sigma}{2p} + \frac{(y - \mu)^2 + \Sigma}{2r} - \tfrac{1}{2} \ln \Sigma + \text{const}$$

Worked example: the Kalman update is free-energy minimisation

Take the prior $s \sim \mathcal{N}(m, p)$, an observation $y \sim \mathcal{N}(s, r)$, and the belief $q = \mathcal{N}(\mu, \Sigma)$. Every term is a Gaussian expectation:

$$F(\mu, \Sigma) = \frac{(\mu - m)^2 + \Sigma}{2p} + \frac{(y - \mu)^2 + \Sigma}{2r} - \tfrac{1}{2} \ln \Sigma + \text{const}$$

$\partial F / \partial \mu = 0$:

$$\mu = \frac{m/p + y/r}{1/p + 1/r}$$

the precision-weighted mean.

$\partial F / \partial \Sigma = 0$:

$$\frac{1}{\Sigma} = \frac{1}{p} + \frac{1}{r}$$

precisions add.

Worked example: the Kalman update is free-energy minimisation

Take the prior $s \sim \mathcal{N}(m, p)$, an observation $y \sim \mathcal{N}(s, r)$, and the belief $q = \mathcal{N}(\mu, \Sigma)$. Every term is a Gaussian expectation:

$$F(\mu, \Sigma) = \frac{(\mu - m)^2 + \Sigma}{2p} + \frac{(y - \mu)^2 + \Sigma}{2r} - \tfrac{1}{2} \ln \Sigma + \text{const}$$

$\partial F / \partial \mu = 0$:

$$\mu = \frac{m/p + y/r}{1/p + 1/r}$$

the precision-weighted mean.

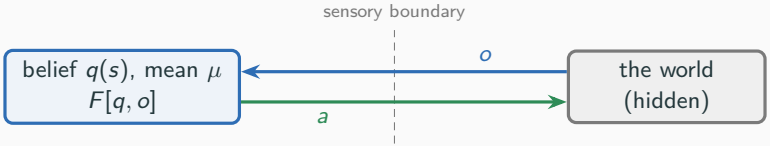
$\partial F / \partial \Sigma = 0$:

$$\frac{1}{\Sigma} = \frac{1}{p} + \frac{1}{r}$$

precisions add.

The two halves come from different terms. The *mean* comes from the prediction errors; the *variance* comes from the curvature of F balanced against the entropy $-\frac{1}{2} \ln \Sigma$. Regularised least squares is the same trade with a flat prior.

One quantity, differentiated twice



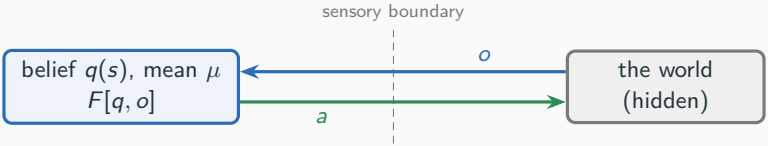
Perception: fit the belief to the data

$$\dot{\mu} = -\frac{\partial F}{\partial \mu}$$

Action: fit the data to the belief

$$\dot{a} = -\frac{\partial F}{\partial a} = -\frac{\partial F}{\partial o}\frac{\partial o}{\partial a}$$

One quantity, differentiated twice



Perception: fit the belief to the data

$$\dot{\mu} = -\frac{\partial F}{\partial \mu}$$

Action: fit the data to the belief

$$\dot{a} = -\frac{\partial F}{\partial a} = -\frac{\partial F}{\partial o}\frac{\partial o}{\partial a}$$

$\partial o/\partial a$ is the causal model running forwards, and a scan already in the buffer cannot be changed. So a decision scores the observations it *would* bring:

$$G(a) = \mathbb{E}_{q(o,s \mid a)}\big[F\big] \approx \underbrace{-I(s; o \mid a)}_{\text{epistemic: what I would learn}} \underbrace{-\mathbb{E}_{q(o \mid a)} \ln p(o)}_{\text{pragmatic: what I want to see}}$$

source: active_inference.tex l.335–427; derived in full later in the week

A choice, marked as a choice

- Closeness was measured with KL. Nothing forced that: it is the $\alpha \to 1$ member of the Rényi family, and the bound survives the substitution.
- The order α controls the *character* of the error: does the belief err towards too broad, or towards too confident?
- On this robot, that is not a statistical nicety. A covariance is spent: on deciding where to look next, and on how far a measurement may move a belief.

The thesis keeps $\alpha \to 1$ throughout. Making α itself a latent variable is open future work.

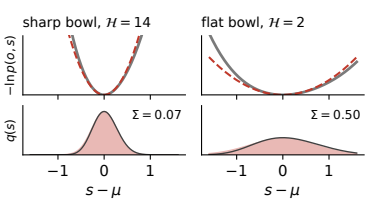
source: active_inference.tex l.270–286; future_work.tex l.22

The Laplace approximation

Gaussian beliefs: the form everything later takes

Take $q(s) = \mathcal{N}(\mu, \Sigma)$ and expand around the mode:

$$F \approx -\ln p(o, \mu) + \tfrac{1}{2}\,\mathrm{tr}(\mathcal{H}\Sigma) - \tfrac{1}{2}\ln|\Sigma| - \tfrac{n}{2}\ln(2\pi e), \qquad \mathcal{H} = -\nabla_s^2 \ln p(o, s)\big|_\mu$$



The **mean** moves to the bottom of the bowl: nonlinear least squares.

The **covariance** is pulled two ways, curvature against entropy; $\partial F / \partial \Sigma = 0$ balances them.

dashed: the parabola Laplace sees; thin: the exact density.

$\boxed{\Sigma = \mathcal{H}^{-1}}$ The uncertainty comes as a by-product of fitting, at no extra modelling cost.

source: active_inference.tex l.310–333

Gauss–Newton, in one line

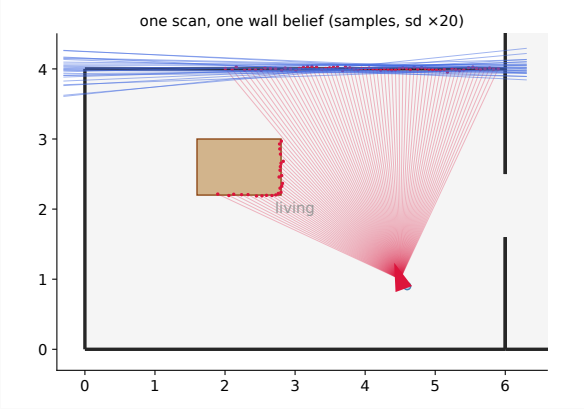
For whitened residuals $r(s)$ with noise covariance R, and a Gaussian prior $\mathcal{N}(m, P)$:

$$-\ln p(o, s) = \tfrac{1}{2} r^\top R^{-1} r + \tfrac{1}{2}(s - m)^\top P^{-1}(s - m) + \text{const},$$

$$\mathcal{H} \approx J^\top R^{-1} J + P^{-1}, \qquad \Delta s = -\mathcal{H}^{-1}\big(J^\top R^{-1} r + P^{-1}(s - m)\big).$$

- Iterate to convergence. The final $\mathcal{H}^{-1}$ *is* the belief's covariance.
- In the lab, `laplace(residual, x0, sigma, prior=...)` does exactly this. It also returns the Laplace free energy $F \approx -\ln p(o, \mu) + \frac{1}{2}\ln|\mathcal{H}| - \frac{n}{2}\ln 2\pi$, which you will use to compare models.

Example: one wall as a belief



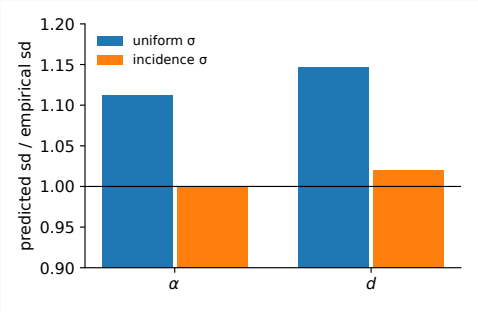
A line with normal $n(\alpha)$ at distance d:

$$r_i = \cos\alpha\, x_i + \sin\alpha\, y_i - d$$

$$\frac{\partial r_i}{\partial(\alpha, d)} = \left[\, -\sin\alpha\, x_i + \cos\alpha\, y_i,\ -1 \,\right]$$

- The first column is a **lever arm**. Points far apart along the wall pin the angle.
- α and d come out correlated ($\rho \approx -0.97$), because d is measured from an origin the points are far away from.

Is the covariance honest?



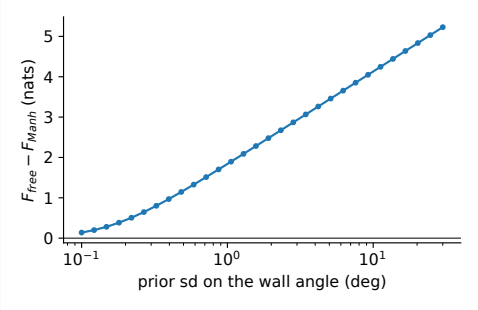
$\Sigma = \mathcal{H}^{-1}$ is a *claim* about how far the fit would move on a re-scan. Test it with 400 scans from one pose:

- uniform σ: about 10% too pessimistic, and it does not average away;
- the noise is along the *beam*, the residual along the wall's *normal*: $\sigma_i = \sigma\,|u_i^\top n|$;
- with that likelihood: within a few per cent.

A Laplace covariance is exactly as honest as the likelihood you wrote.

source: lab 1, exercises 3–4; the real localiser's incidence weight: room_concept.tex l.232

Free energy compares models: the price of a parameter



Two models of the same wall:

- **Manhattan**: the angle is fixed at 90°; one parameter, d.
- **free**: $\alpha \sim \mathcal{N}(90^\circ, s^2)$; two parameters.

Both fit equally well. The free model pays $\approx \ln(s/\sigma_\alpha)$ nats of *complexity*: the cost of squeezing the prior down to the posterior.

Tilt the points by 1.5° and the verdict flips by 70 nats.

Nobody wrote "if the angle error exceeds k degrees". That is the first instance of **no thresholds**.

What measurement said

Free energy is not computation

Measured: a better model does not save arithmetic

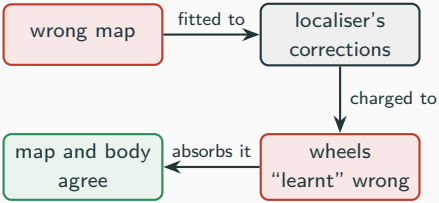
Injecting a body-model error from 0 to 10%:

- the solver iterations per correction stay flat to within a twentieth;
- the size of the correction grows by about half as much again.

- The agent pays in *how far its beliefs must move*. The arithmetic doing the moving is nearly constant.
- "A better model saves the agent work" has to say *which* work. The two readings come apart, and an endpoint chosen on the wrong reading measures nothing.

source: active_inference.tex l.835–848

A factor inherits the errors of what it is estimated through



It gets *quieter*, not louder: internally consistent, externally wrong.

Measured: an omitted parameter

Admitting a camera-mount translation the data cannot support moves the reported pitch by **nine times** its stated interval. An interval speaks only about the parameters the model contains.

Carry an external reference. Never trust only what the agent computes about itself.

source: active_inference.tex l.899–950

Lab

Lab 1: probability by hand, then a wall as a belief

`labs/day1.ipynb`, about 50 minutes. Today the robot knows its pose and we cheat on which points are the wall's.

1. **Warm-up**: the door table by hand (marginalise, condition, both factorisations); read an anti-causal order off a joint; condition a Gaussian in both forms.
2. Write the wall's residual and Jacobian; fit it with `laplace`.
3. **Honesty test**: 500 scans; predicted vs empirical σ. Find the 10% and fix the likelihood.
4. **Model choice**: Manhattan vs free, by ΔF. Then tilt the wall.

Where to look next waits for Day 4, when G is derived in full.

Tomorrow

Today there was one robot, one belief and one F. A real robot runs a dozen processes:

- where does the belief live, and who holds it?
- how is one free energy divided among processes that never talk to each other?

And the commitment that runs through the rest of the week: a gate is what a model looks like before the modelling has been done.

Reading for Day 2

Thesis chapters:

- *The CORTEX Architecture*, §§"The working memory as a distributed graph", "Active Inference in CORTEX";
- *Concept Agents*, up to "A design rule that follows from the bound".